

# Neural Networks

Introduction to Data Science  
Spring 1404

Yadollah Yaghoobzadeh

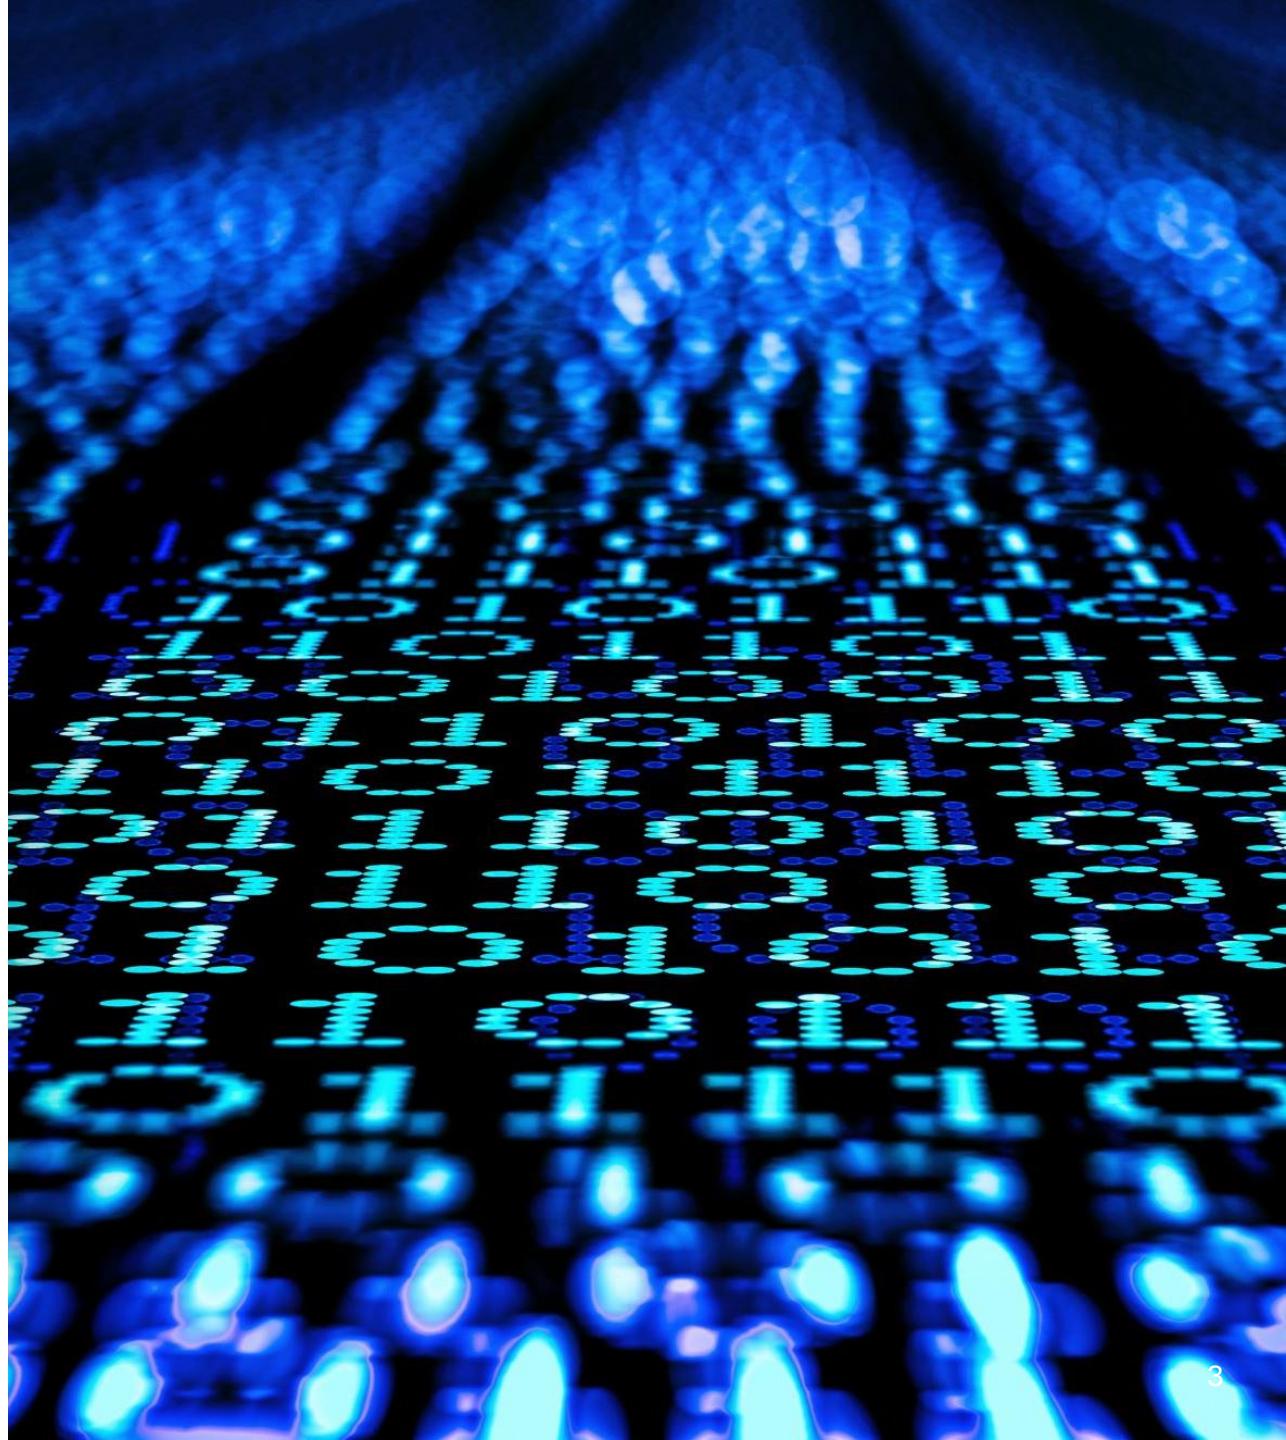
# Deep Learning

Multi-layer artificial neural networks



# Rise of DL

- Data
- Compute
- Algorithms



# Deep learning

- What is its position wrt to ML and AI?

# Key differences with other ML techniques

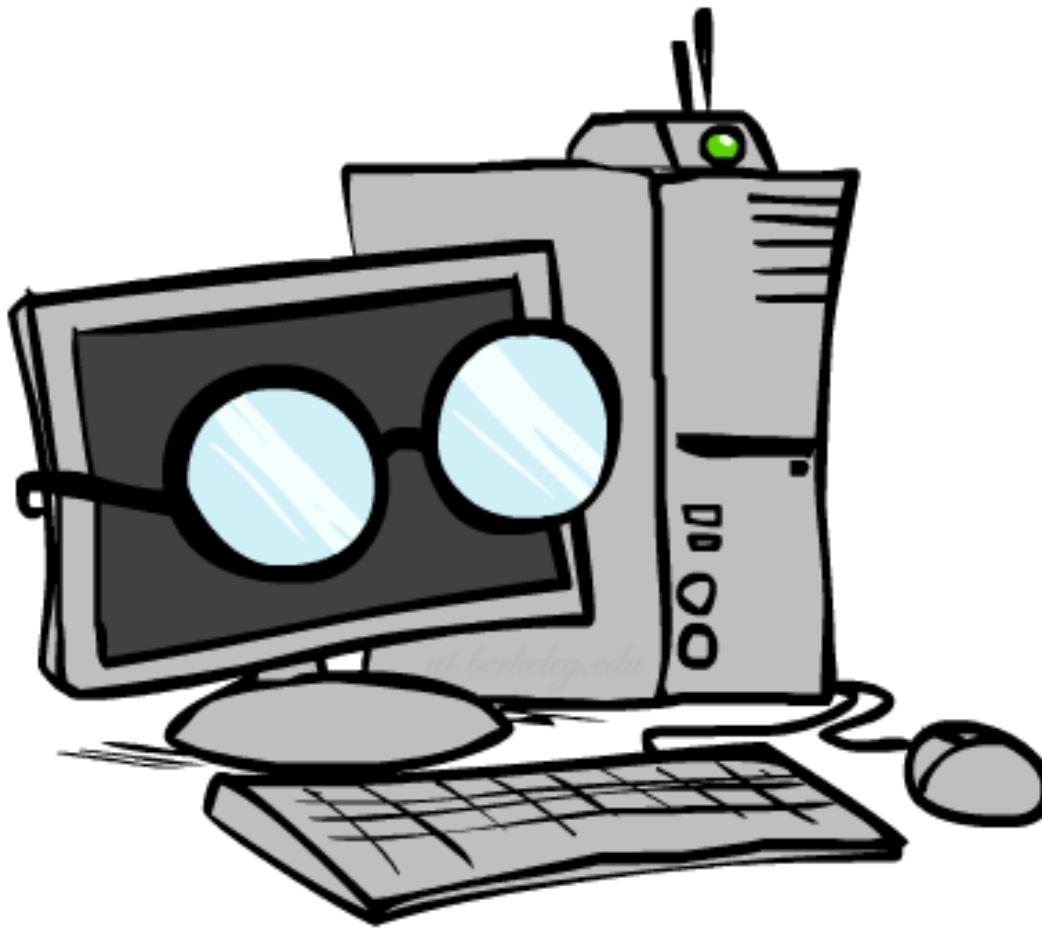
- It learns the features itself
- It is scalable and learn from any large data

# Applications

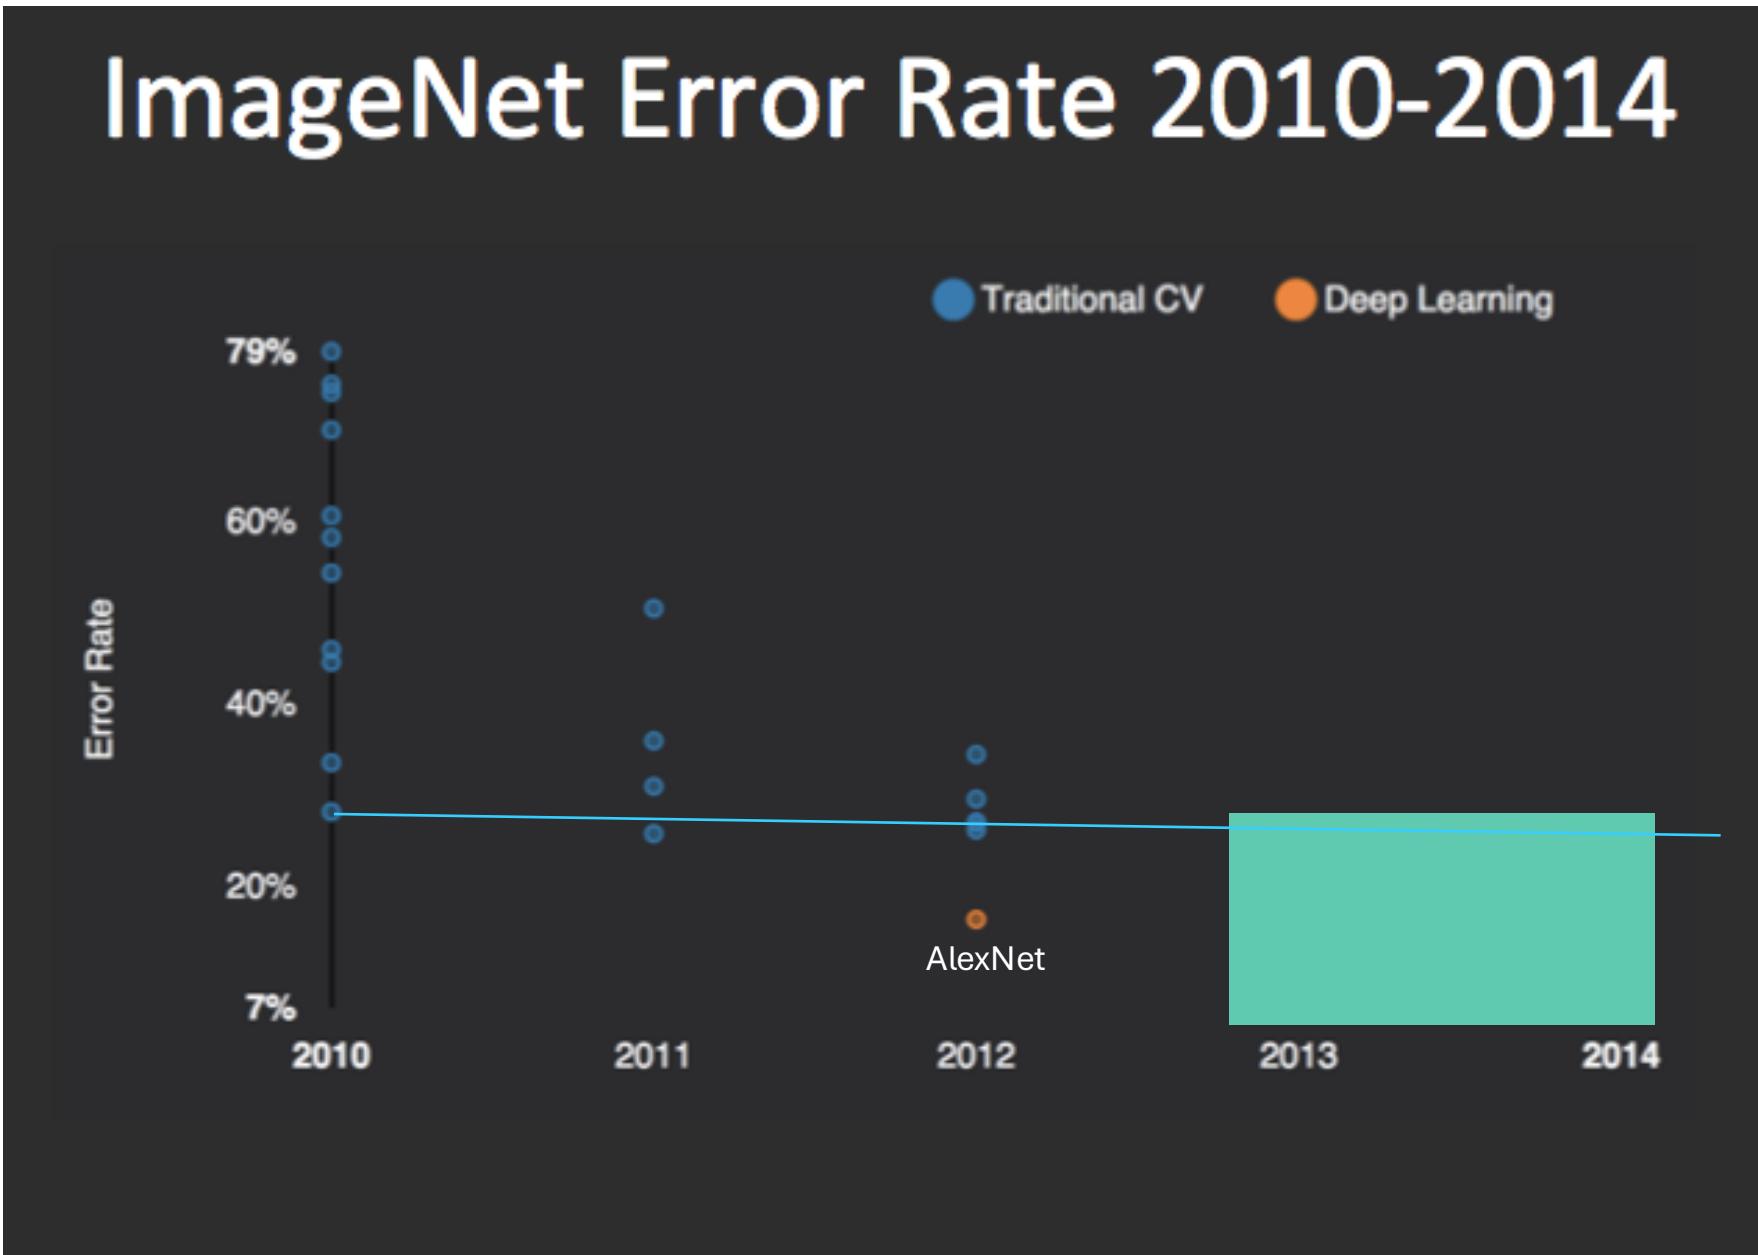
- Natural language processing
  - Machine translation
  - Search
  - Sentiment analysis
  - ➔ Large language models
- Computer vision
  - Image captioning
  - Image generation
  - Image classification
  - Image segmentation
- Recommendation systems

# How well does it work?

# Computer Vision

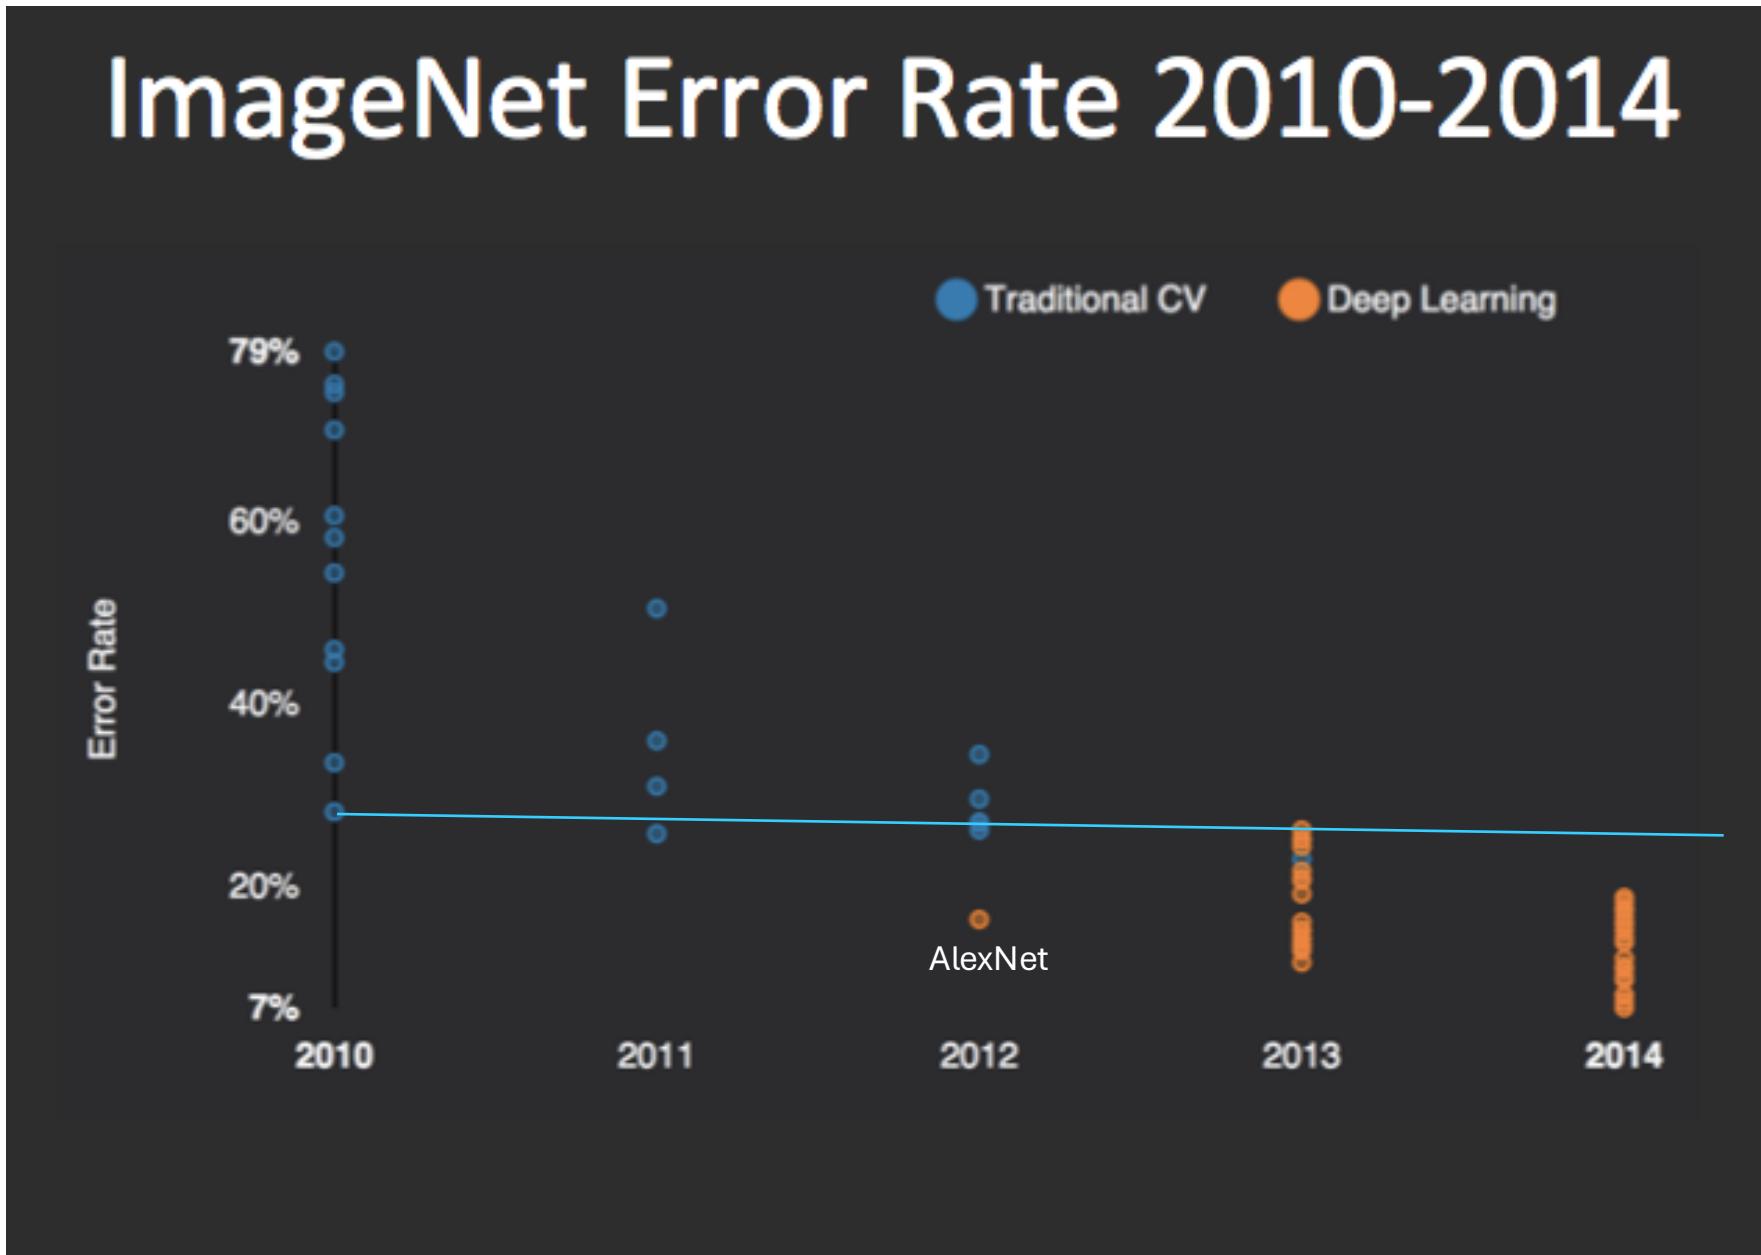


# Performance

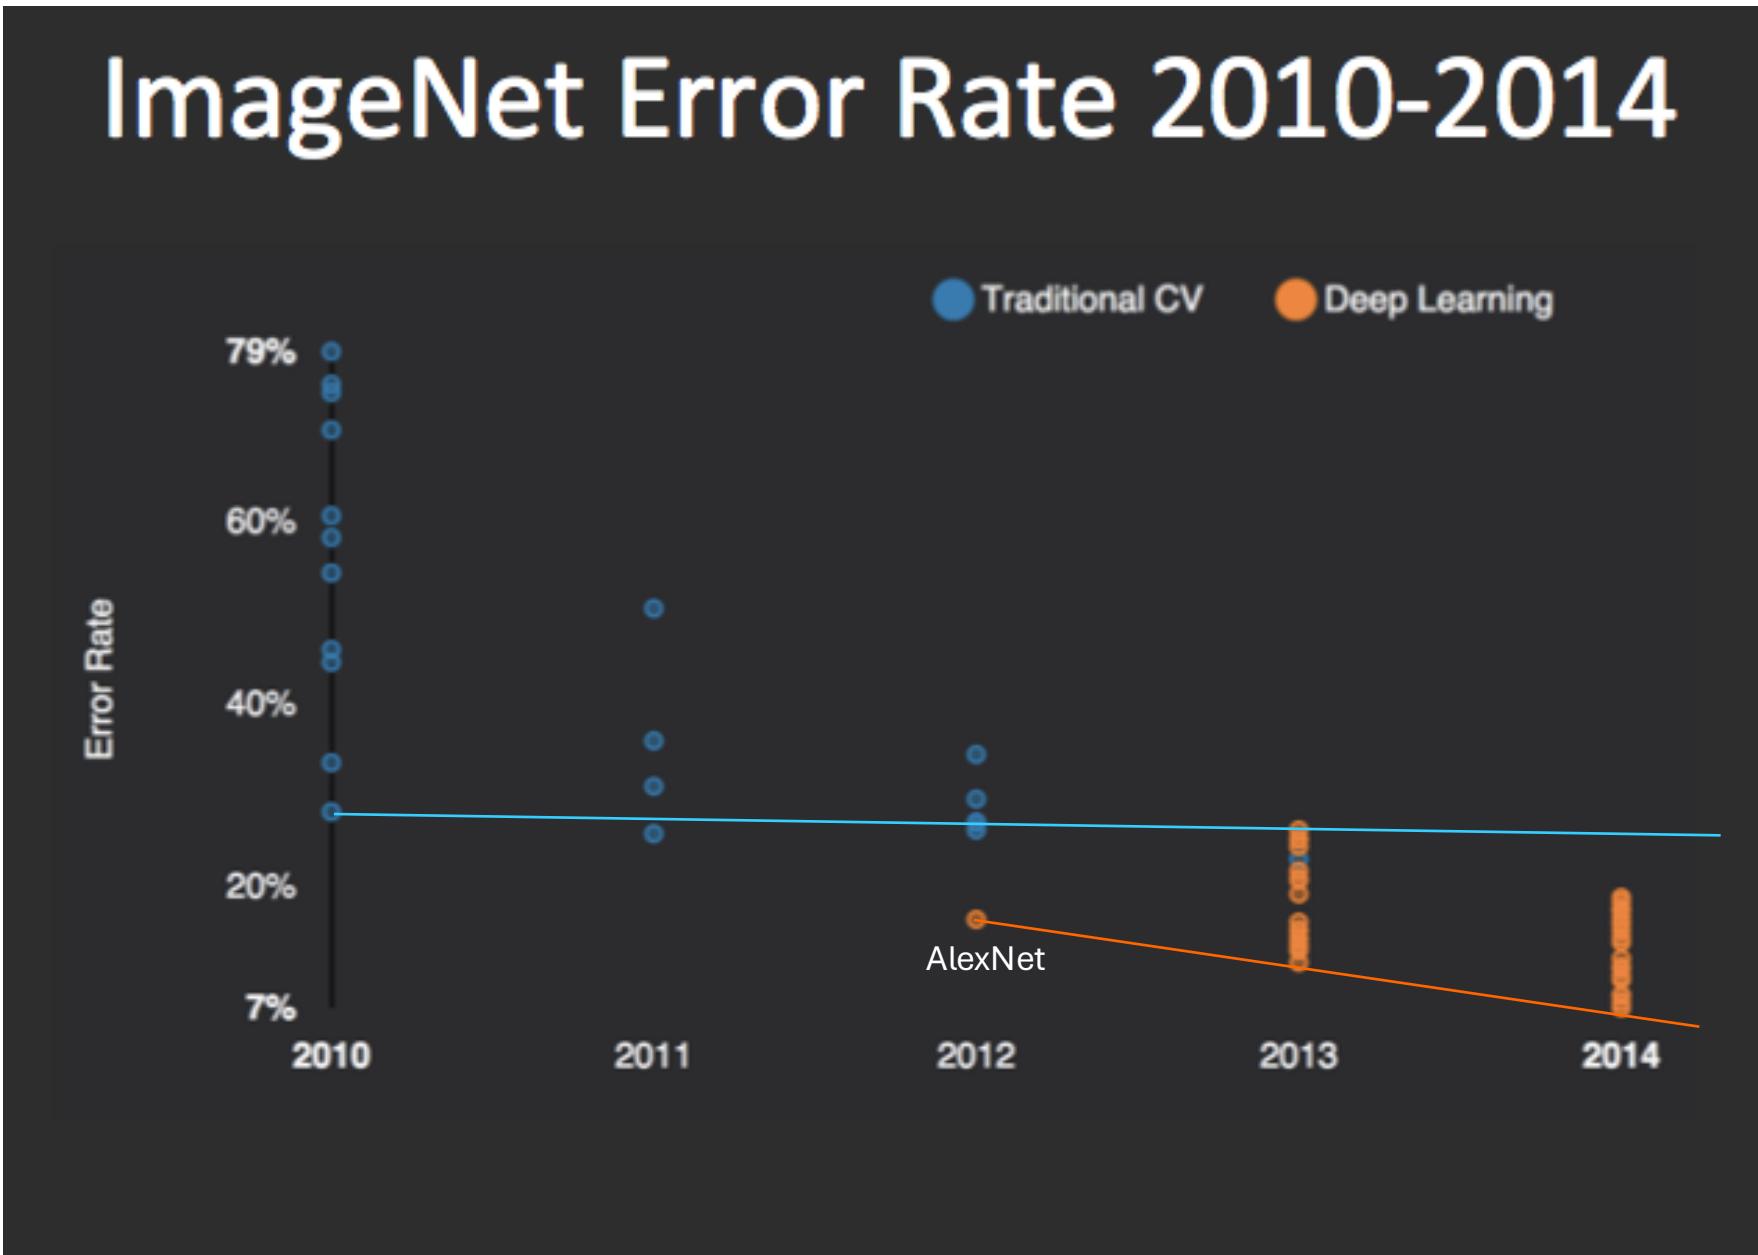


graph credit Matt  
Zeiler, Clarifai

# Performance

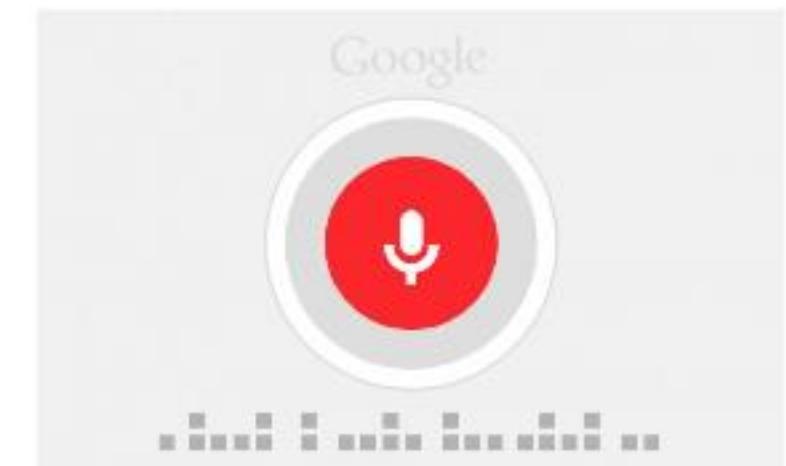
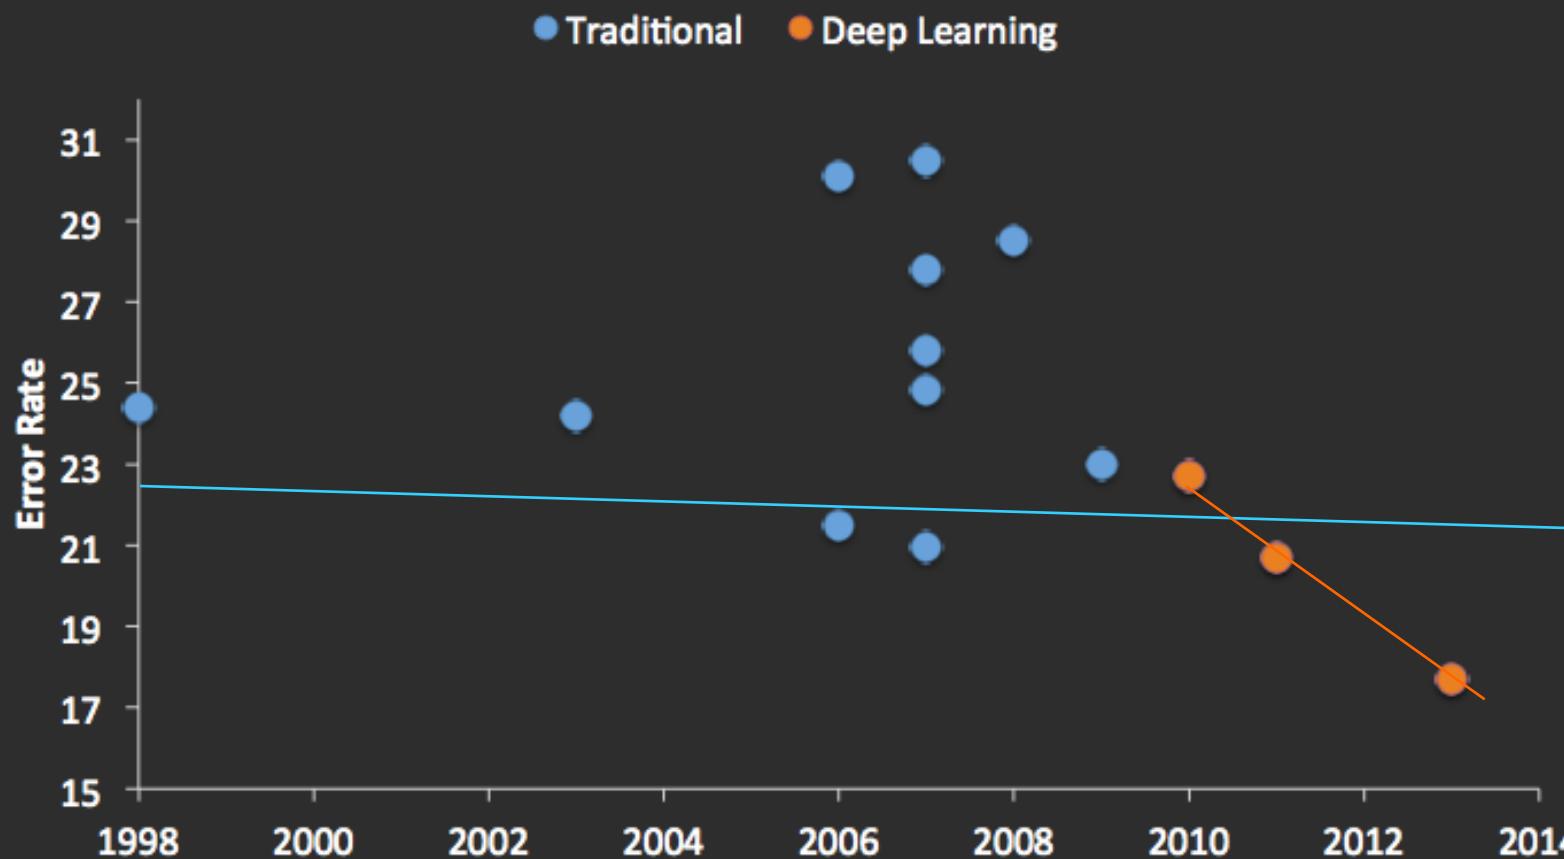


# Performance



# Speech Recognition

## TIMIT Speech Recognition



graph credit Matt Zeiler, Clarifai

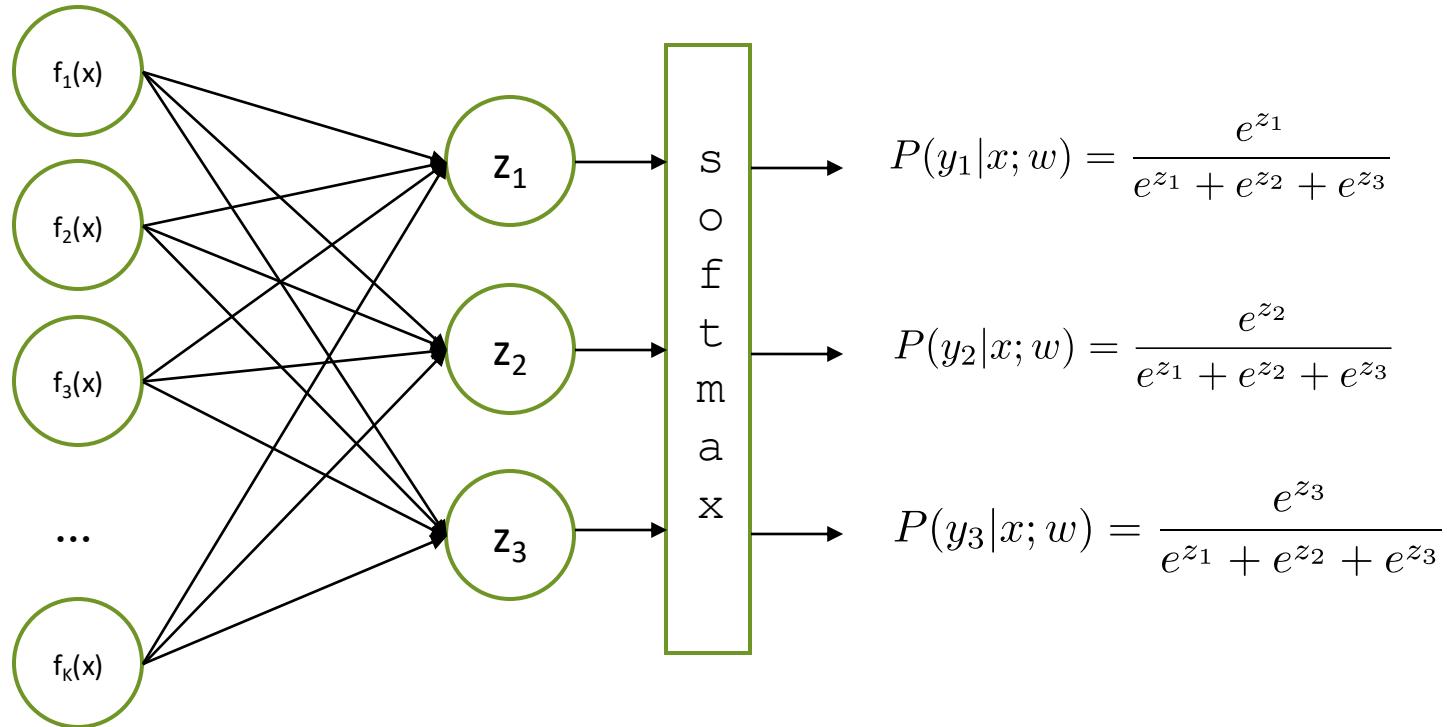
# Natural language processing

- BERT (2017)
- ChatGPT (2022)

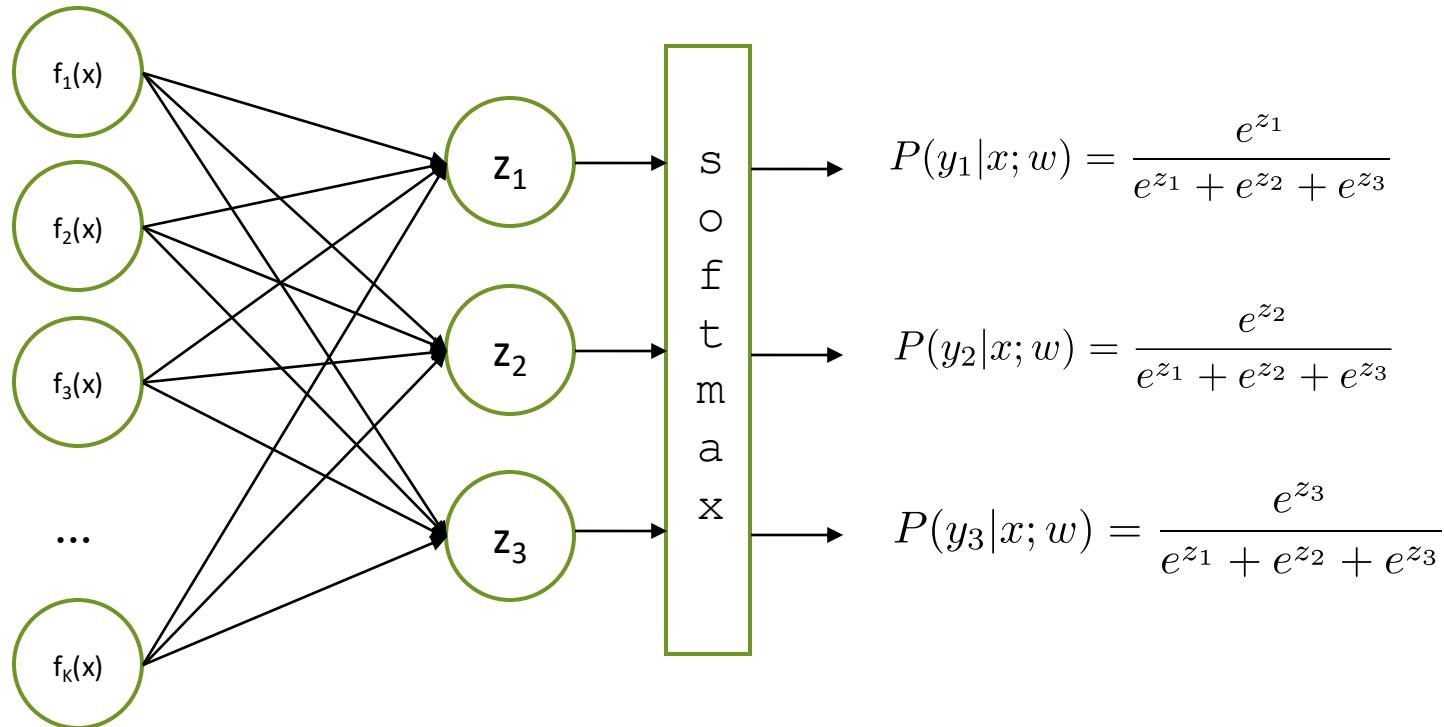
# But what is ANN?

# Multi-class Logistic Regression

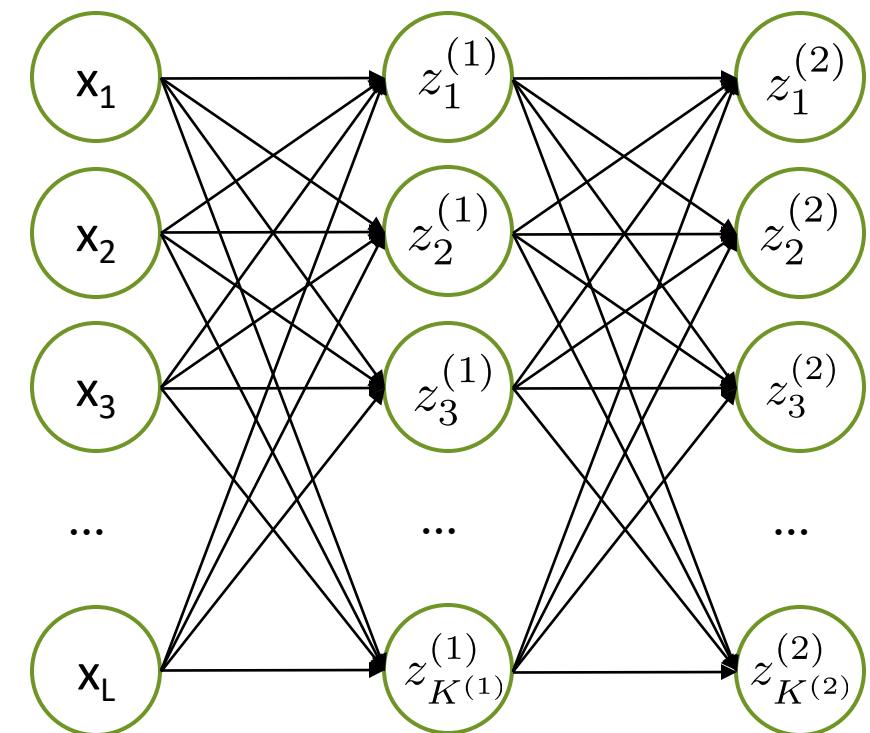
= special case of neural network



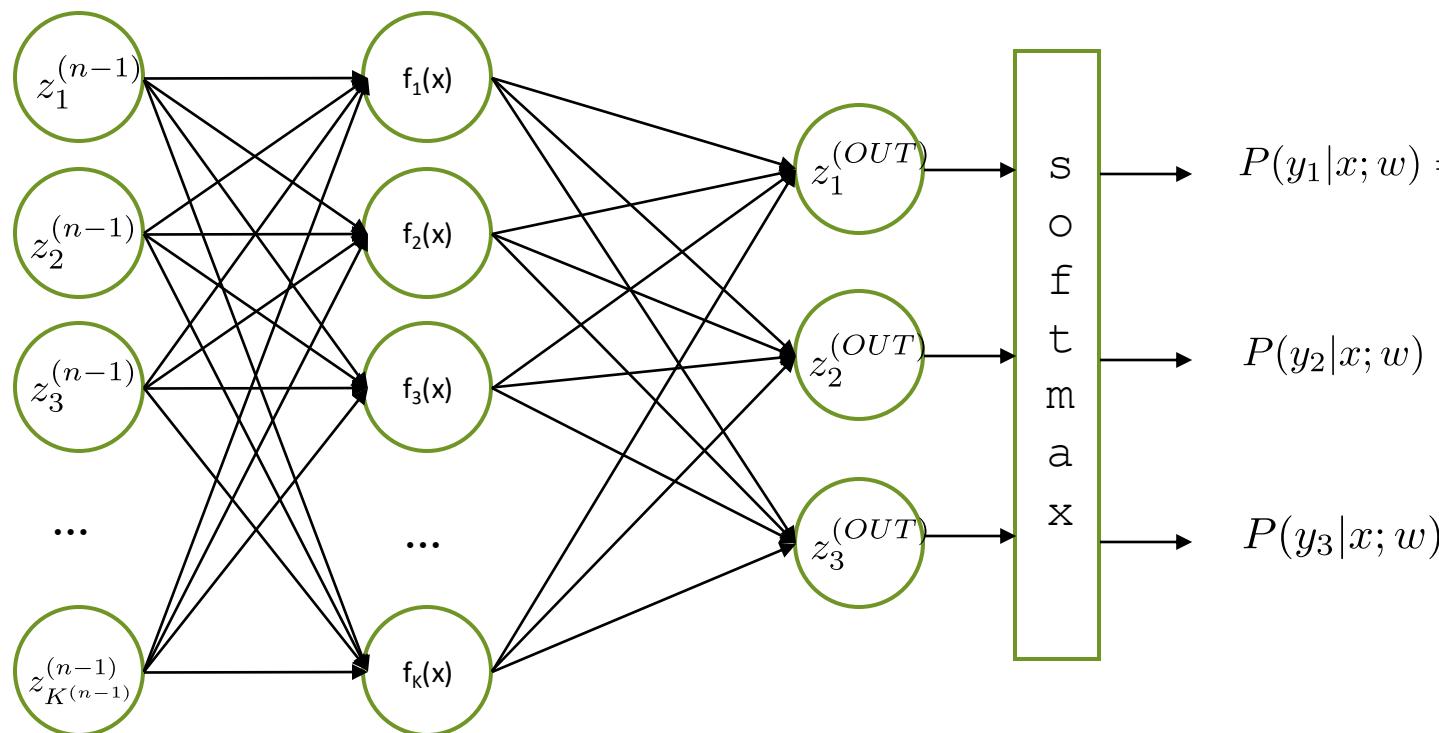
# Deep Neural Network = Also learn the features!



# Deep Neural Network = Also learn the features!



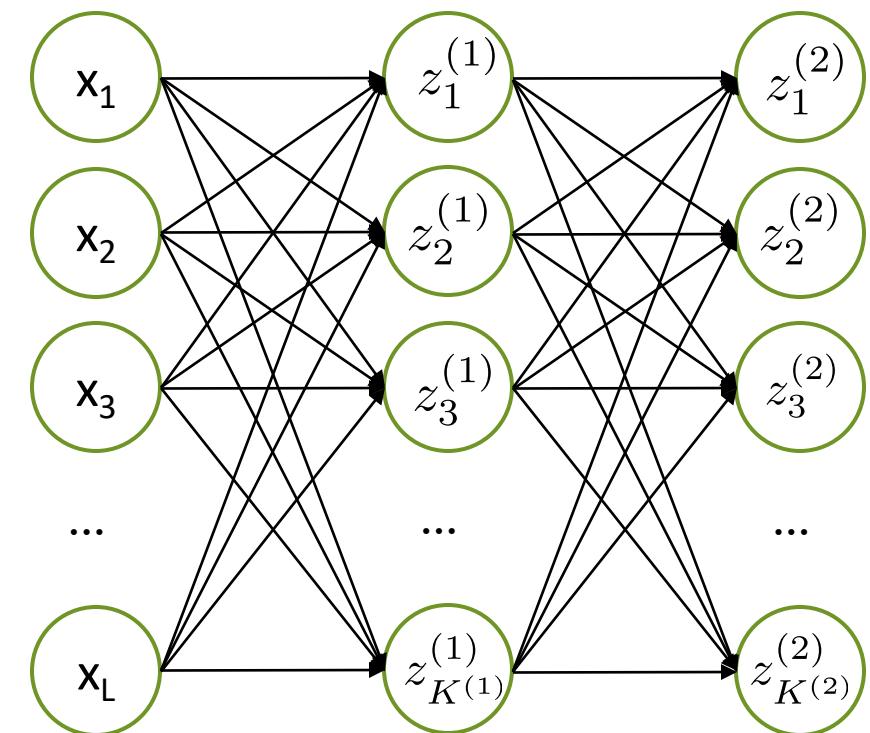
...



$$z_i^{(k)} = g\left(\sum_j W_{i,j}^{(k-1,k)} z_j^{(k-1)}\right)$$

**g = nonlinear activation function**

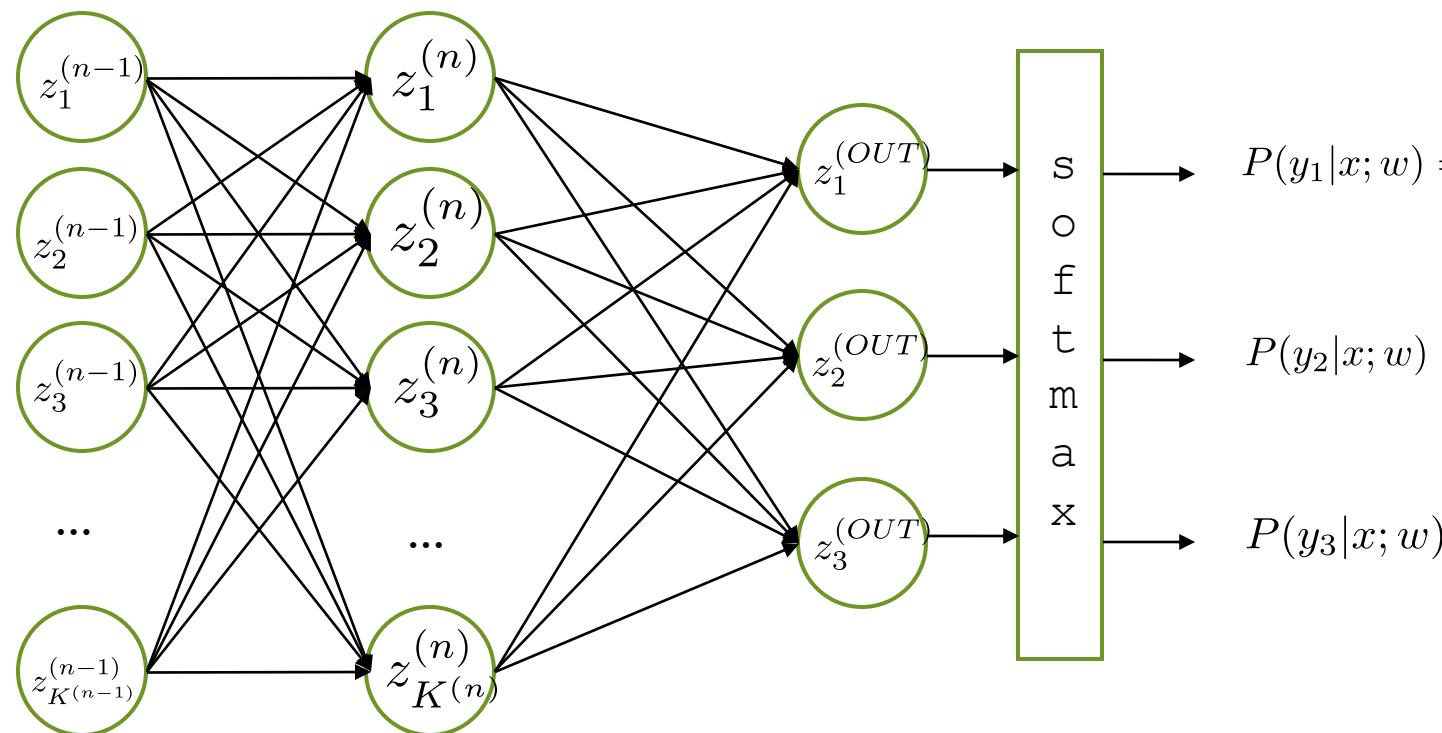
# Deep Neural Network = Also learn the features!



...

...

...

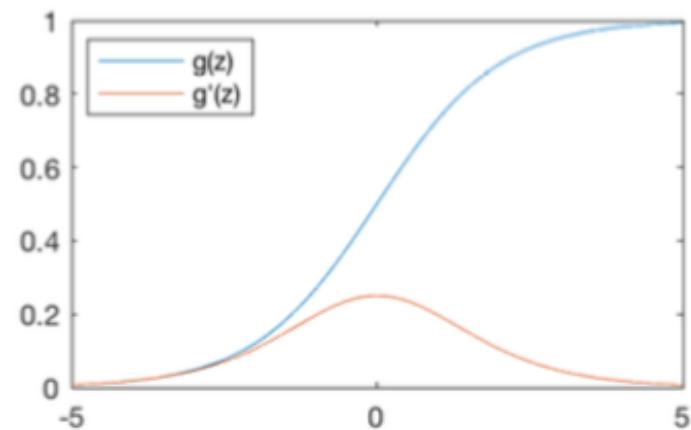


$$z_i^{(k)} = g\left(\sum_j W_{i,j}^{(k-1,k)} z_j^{(k-1)}\right)$$

**g = nonlinear activation function**

# Common Activation Functions

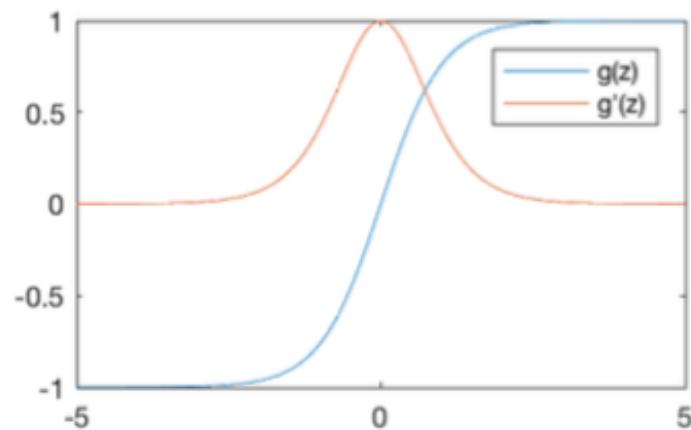
Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

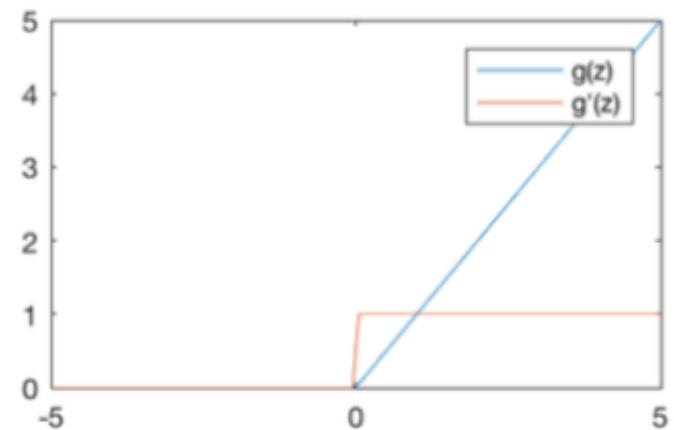
Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

# Deep Neural Network: Also Learn the Features!

Training the deep neural network is just like logistic regression:

$$\max_w \text{ll}(w) = \max_w \sum_i \log P(y^{(i)} | x^{(i)}; w)$$

just  $w$  tends to be a much, much larger vector ☺

→ just run gradient ascent

+ stop when log likelihood of hold-out data starts to decrease

# Optimization Procedure: Gradient Ascent

- init  $w$
- for iter = 1, 2, ...

$$w \leftarrow w + \alpha * \nabla g(w)$$

- $\alpha$ : learning rate --- tweaking parameter that needs to be chosen carefully
- How? Try multiple choices
  - Crude rule of thumb: update changes  $w$  about 0.1 – 1 %

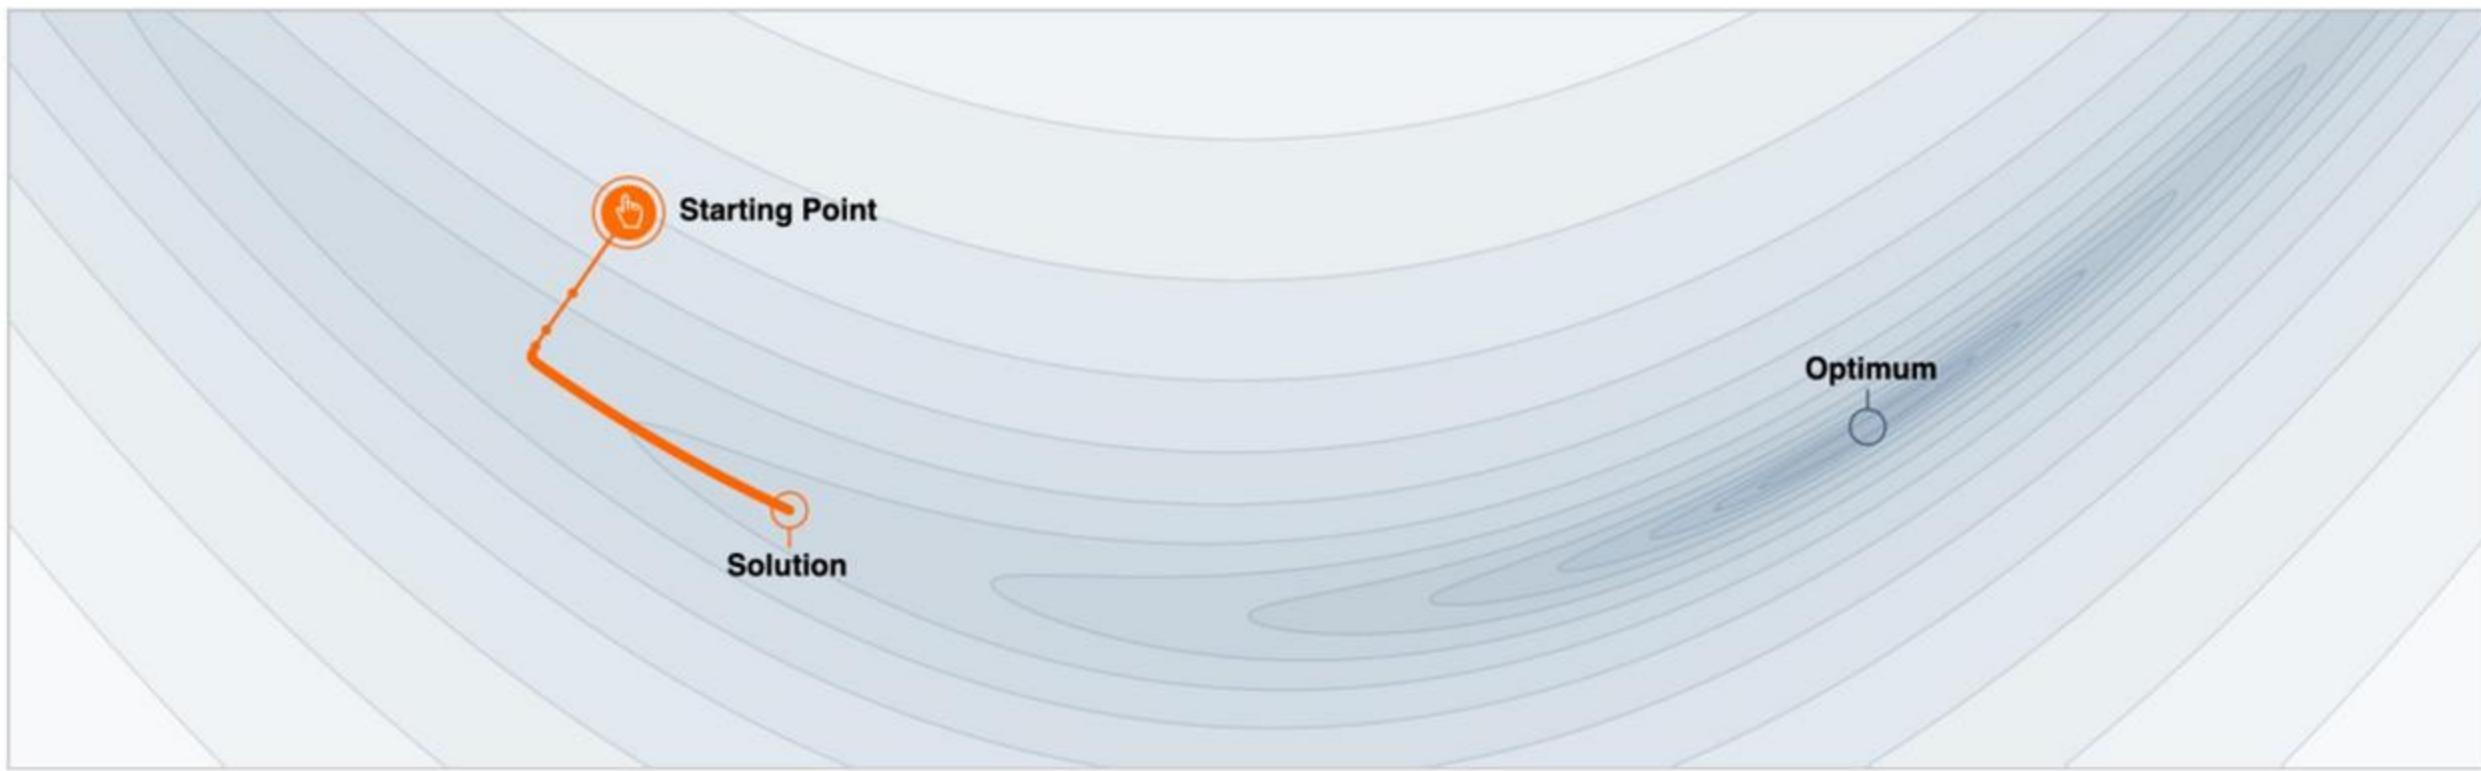
# Minimizing loss functions

- In practice we minimize the loss instead of maximizing the objective function
- E.g.,
  - Cross-entropy loss for classification
  - Mean squared error for regression

# Learning Rate

Choice of learning rate  $\alpha$  is a hyperparameter

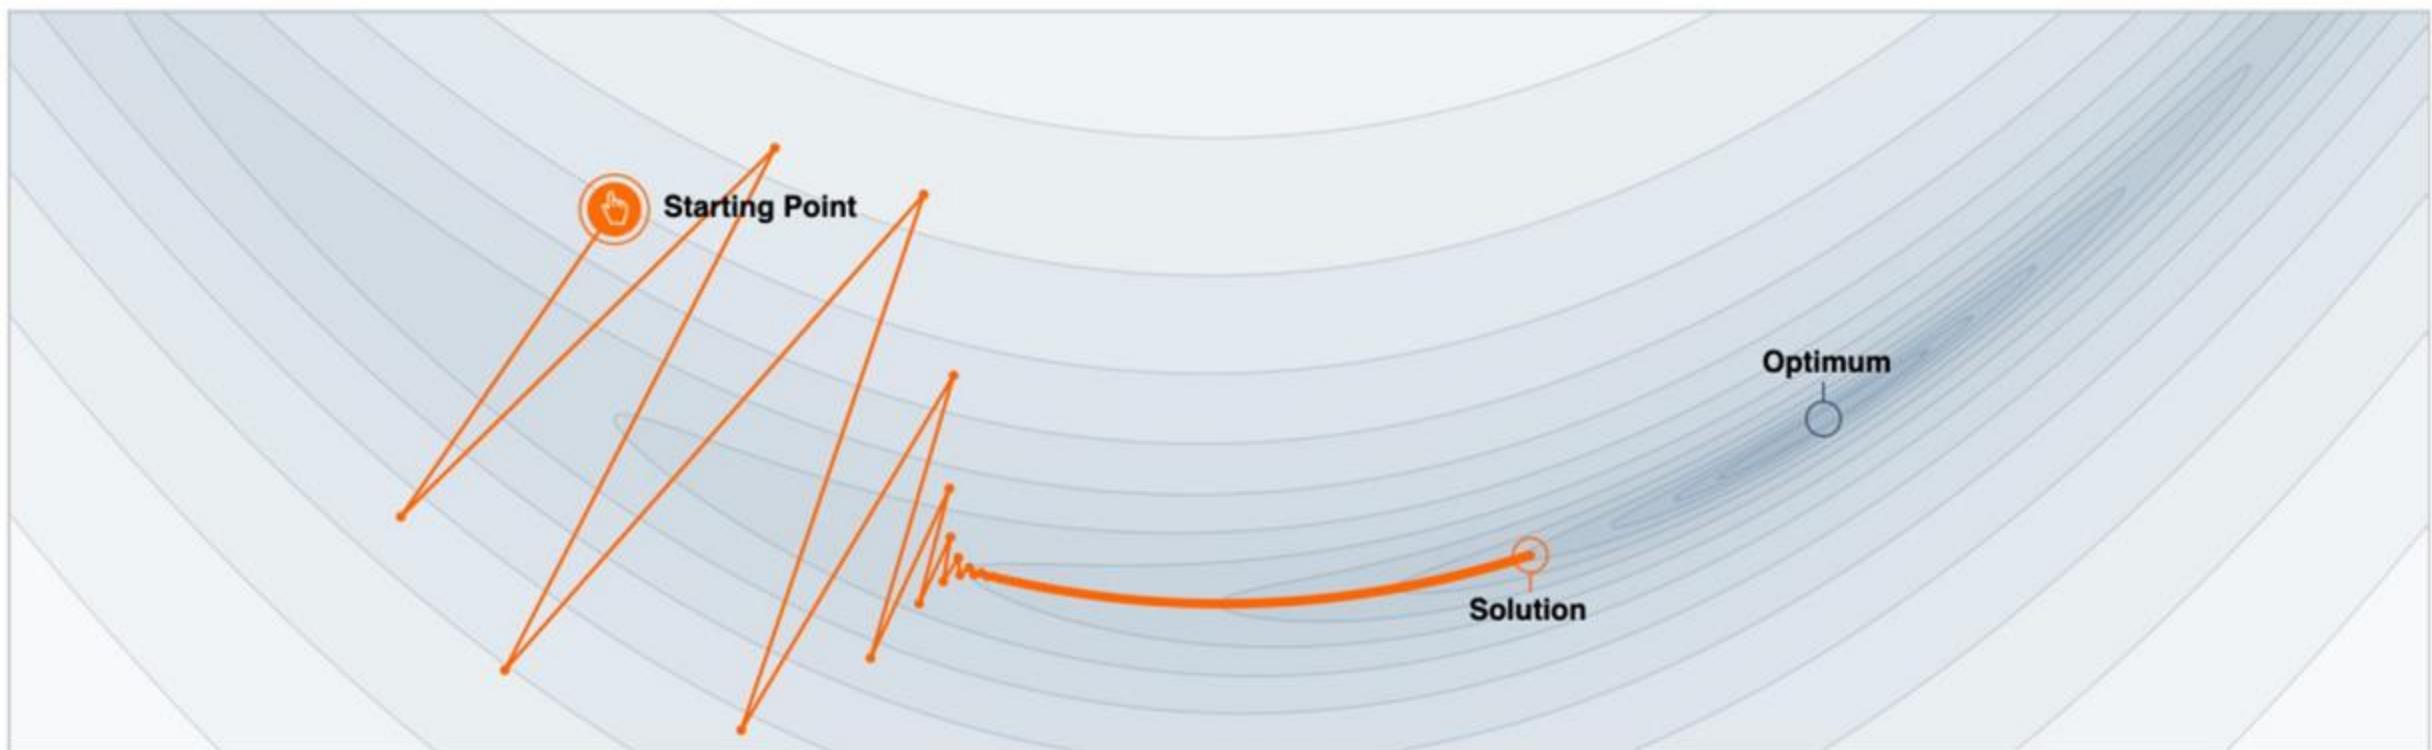
Example:  $\alpha=0.001$  (too small)



# Learning Rate

Choice of step size  $\alpha$  is a hyperparameter

Example:  $\alpha=0.004$  (too large)



# Fun Neural Net Demo Site

- Demo-site:
  - <http://playground.tensorflow.org/>

# Backpropagation

- **Back-propagation:** gradients are computed in the direction from output to input layers and combined using chain rule

$$f(x) = g(h(x))$$

- CHAIN RULE:

If

$$f'(x) = g'(h(x))h'(x)$$

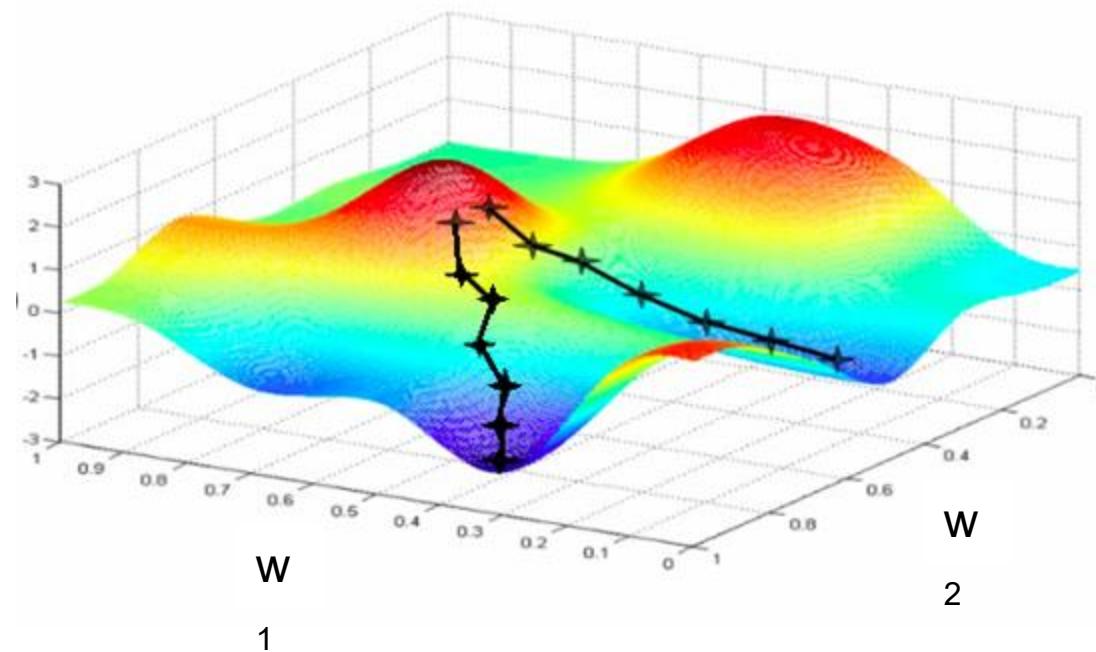
Then

→ Derivatives can be computed by following well-defined procedures

# Deep Neural Network: Training

- Update weights by **gradient descent**:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial E}{\partial \mathbf{w}}$$



# Automatic Differentiation

- Automatic differentiation software
  - e.g. Theano, TensorFlow, PyTorch, Chainer
  - Only need to program the function  $g(x,y,w)$
  - Can automatically compute all derivatives w.r.t. all entries in  $w$
  - This is typically done by caching info during forward computation pass of  $f$ , and then doing a backward pass = “backpropagation”
  - Autodiff / Backpropagation can often be done at computational cost comparable to the forward pass
- Need to know this exists
- How this is done? -- outside of scope of this class

# Summary of Key Ideas

- Optimize probability of label given input 
$$\max_w ll(w) = \max_w \sum_i \log P(y^{(i)}|x^{(i)}; w)$$
- Continuous optimization
  - Gradient ascent:
    - Compute steepest uphill direction = gradient (= just vector of partial derivatives)
    - Take step in the gradient direction
    - Repeat (until held-out data accuracy starts to drop = “early stopping”)
- Deep neural nets
  - Last layer = still logistic regression
  - Now also many more layers before this last layer
    - = computing the features
    - → the features are learned rather than hand-designed
  - Universal function approximation theorem
    - If neural net is large enough
    - Then neural net can represent any continuous mapping from input to output with arbitrary accuracy
    - But remember: need to avoid overfitting / memorizing the training data → early stopping!
  - Automatic differentiation gives the derivatives efficiently (how? = outside of scope of this class)

# The Training Loop

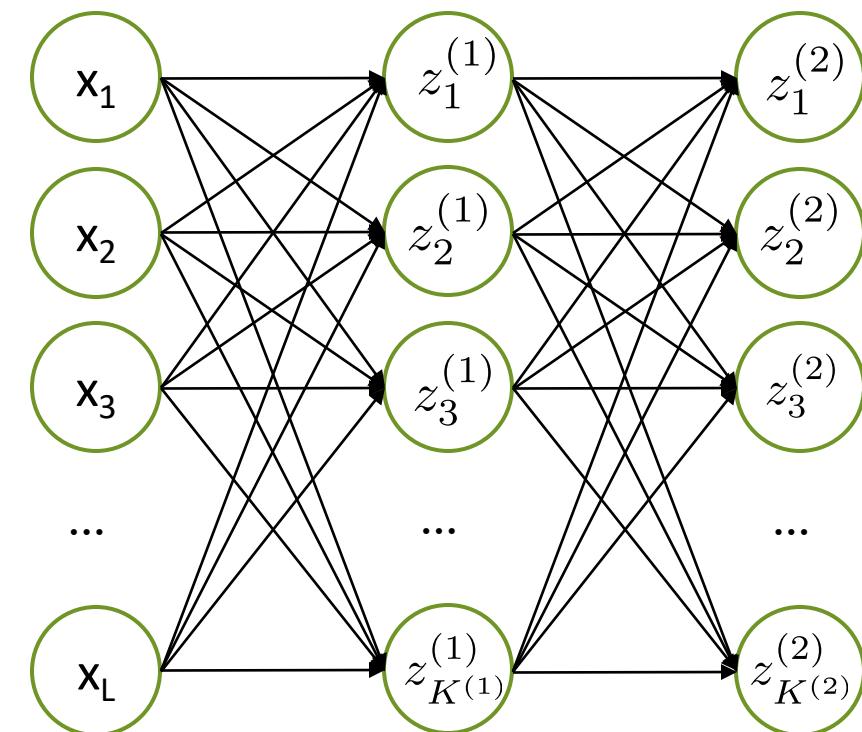
Repeat for multiple epochs

- Get a batch of data
- Forward pass: Calculate predictions
- Calculate Loss
- Backward pass: Calculate gradients (Backprop)
- Update weights/biases (Optimizer)
- We usually have a number of hyperparameters to tune
  - Optimizer hyperparameters like learning rate
  - Batch size
  - Epochs

# Connection Architecture

- Fully connected
- Convolution
- Recurrent
- Self-attention

# Fully connected layers



...

