



SQL – part I

Introduction to Data Science
Spring 1404

Yadollah Yaghoobzadeh

Goals for Today's Lecture

Stepping away from Python and pandas

- Recognizing situations where we need “bigger” tools for manipulating data
- Writing our first database queries

Agenda

- Why Databases?
- Intro to SQL
- Tables and Schema
- Basic Queries
- Grouping

Why Databases?

- **Why Databases?**
- Intro to SQL
- Tables and Schema
- Basic queries
- Grouping

So Far: CSV Files and pandas

- So far, we've worked with data stored in **CSV files**

Berkeley_PD_-_Calls_for_Service.csv

pd.read_csv

The diagram illustrates the process of reading a CSV file into a pandas DataFrame. An arrow points from the text "Berkeley_PD_-_Calls_for_Service.csv" to a box containing the function "pd.read_csv". Another arrow points from this box down to a table representing the resulting DataFrame.

	CASENO	OFFENSE	EVENTDT	EVENTTM	CVLEGEND	CVDOW	InDbDate	Block_Location	BLKADDR	City	State
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:58	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\\n(37.869058, -122.270455)	NaN	Berkeley	CA
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:38	LARCENY	4	06/15/2021 12:00:00 AM	Berkeley, CA\\n(37.869058, -122.270455)	NaN	Berkeley	CA
2	21090494	THEFT MISD. (UNDER \$950)	04/19/2021 12:00:00 AM	12:15	LARCENY	1	06/15/2021 12:00:00 AM	2100 BLOCK HASTE ST\\nBerkeley, CA\\n(37.864908,...)	2100 BLOCK HASTE ST	Berkeley	CA
3	21090204	THEFT FELONY (OVER \$950)	02/13/2021 12:00:00 AM	17:00	LARCENY	6	06/15/2021 12:00:00 AM	2600 BLOCK WARRING ST\\nBerkeley, CA\\n(37.86393...	2600 BLOCK WARRING ST	Berkeley	CA
4	21090179	BURGLARY AUTO	02/08/2021 12:00:00 AM	6:20	BURGLARY - VEHICLE	1	06/15/2021 12:00:00 AM	2700 BLOCK GARBER ST\\nBerkeley, CA\\n(37.86066,...)	2700 BLOCK GARBER ST	Berkeley	CA

- Perfectly reasonable workflow for small data that we're not actively sharing with others.

Brief Databases Overview

A **database** is an organized collection of data.

A **Database Management System (DBMS)** is a software system that **stores, manages, and facilitates access** to one or more databases.



Advantages of DBMS over CSV (or Similar)

Data Storage:

- Reliable storage to survive system crashes and disk failures.
- Optimize to compute on data that does not fit in memory.

Data Management:

- Configure how data is organized and who has access.
- Can enforce guarantees on the data (e.g. non-negative person weight or age).
 - Can be used to prevent data anomalies.
 - Ensures safe concurrent operations on data (multiple users reading and writing simultaneously, e.g. ATM transactions).

Intro to SQL

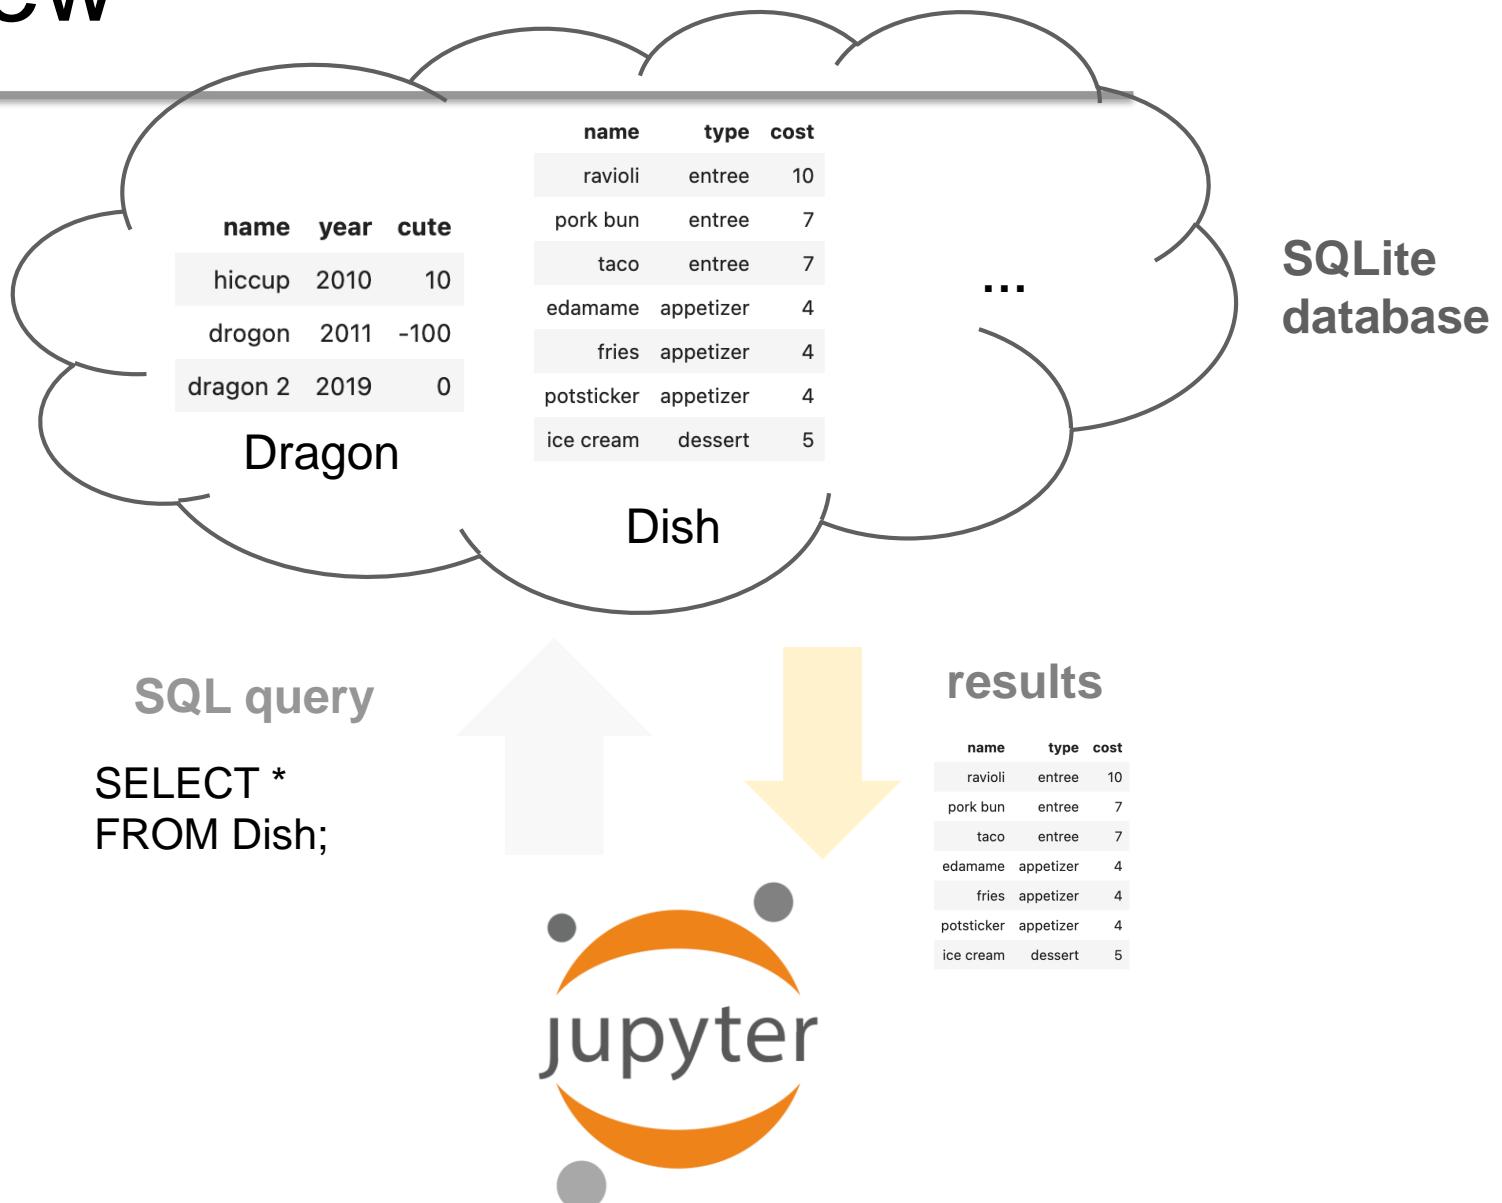
- Why Databases?
- **Intro to SQL**
- Tables and Schema
- Basic queries
- Grouping

SQL

- Today we'll be using a programming language called "Structured Query Language" or SQL.
- SQL is its own programming language, totally distinct from Python.
- SQL is a special purpose programming language used specifically for communicating with databases.
- We will program in SQL using Jupyter notebooks.

How to pronounce? An ongoing debate.

Quick SQL Overview



Tables and Schema

- Why Databases?
- Intro to SQL
- **Tables and Schema**
- Basic queries
- Grouping

SQL Terminology

Column or Attribute or Field

Row or Record or Tuple

name TEXT, PK	year INT, >=2000	cute INT
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0

Dragon → table name

Every column in a SQL table has three properties: **ColName**, **Type**, and zero or more **Constraints**.

(Contrast with pandas: Series have names and types, but no constraints.)

Table Schema

A **schema** describes the logical structure of a table. Whenever a new table is created, the creator must declare its schema.

For each column, specify the:

- **Column name**
- **Data type**
- **Constraint(s) on values**

```
CREATE TABLE Dragon (
    name TEXT PRIMARY KEY,
    year INTEGER CHECK (year >= 2000),
    cute INTEGER
)
```

Repeat for all tables in the database:

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))

Example Types

Some examples of SQL types:

- INT: Integers.
- FLOAT: Floating point numbers.
- TEXT: Strings of text.
- BLOB: Arbitrary data, e.g. songs, video files, etc.
- DATETIME: A date and time.

Note: Different implementations of SQL support different types.

- SQLite: <https://www.sqlite.org/datatype3.html>
- MySQL: <https://dev.mysql.com/doc/refman/8.0/en/data-types.html>

Example Constraints

Some examples of **constraints**:

- **CHECK**: data must obey the given check constraint.
- **PRIMARY KEY**: specifies that this key is used to uniquely identify rows in the table.
- **NOT NULL**: null data
- **DEFAULT**: provides a insertion.

What is this primary key constraint?

type	name	tbl_name	rootpage	sql
table	sqlite_sequence	sqlite_sequence	7	CREATE TABLE sqlite_sequence(name,seq)
table	Dragon	Dragon	2	CREATE TABLE Dragon (name TEXT PRIMARY KEY, year INTEGER CHECK (year >= 2000), cute INTEGER)
table	Dish	Dish	4	CREATE TABLE Dish (name TEXT PRIMARY KEY, type TEXT, cost INTEGER CHECK (cost >= 0))
table	Scene	Scene	6	CREATE TABLE Scene (id INTEGER PRIMARY KEY AUTOINCREMENT, biome TEXT NOT NULL, city TEXT NOT NULL, visitors INTEGER CHECK (visitors >= 0), created_at DATETIME DEFAULT (DATETIME('now')))

Primary Keys

A **primary key** is the set of column(s) used to uniquely identify each record in the table.

- In the Dragon table, the “name” of each Dragon is the primary key.
- In other words, no two dragons can have the same name!
- Primary key is used **under the hood** for all sorts of optimizations.

name TEXT, PK	year INT, ≥ 2000	cute INT
hiccup	2010	10
drogon	2011	-100
Dragon 2	2019	0

Why specify primary keys?
More next time when we discuss JOINs...

Basic Queries

- Why Databases?
- Intro to SQL
- Tables and Schema
- **Basic Queries**
- Grouping

Query Syntax So Far

```
SELECT <column list>  
FROM <table>;
```



; Marks the end of a SQL statement.

New keywords

```
SELECT <column list>  
FROM <table>  
[WHERE <predicate>]  
[ORDER BY <column list>]  
[LIMIT <number of rows>]  
[OFFSET <number of rows>];
```

By the end of this section, you will learn these new keywords!

But first, more SELECT

- Recall our simplest query, which returns the full relation:

```
SELECT *
FROM Dragon;
```

- SELECT specifies the column(s) that we wish to appear in the output. FROM specifies the database table from which to select data.
- *Every* query must include a SELECT clause (how else would we know what to return?) and a FROM clause (how else would we know where to get the data?)
- An asterisk (*) is shorthand for “all columns”.

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None

But first, more SELECT

Recall our simplest query, which returns the full relation:

```
SELECT *  
FROM Dragon;
```

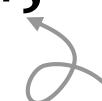


table name

name	year	cute
hiccup	2010	10
drogon	2011	-100
dragon 2	2019	0
puff	2010	100
smaug	2011	None

We can also SELECT only a **subset of the columns**:

column expression list

SELECT cute, year
FROM Dragon;

cute	year
10	2010
-100	2011
0	2019
100	2010
None	2011

Columns selected in specified order

Aliasing with AS

To rename a SELECTed column, use the AS keyword

```
SELECT cute AS cuteness,  
       year AS birth  
FROM Dragon;
```

An **alias** is a name given to a column or table by a programmer. Here, “cuteness” is an alias of the original “cute” column (and “birth” is an alias of “year”)

	cuteness	birth
	10	2010
	-100	2011
	0	2019
	100	2010
	None	2011

SQL Style: Newline Separators

The following two queries both retrieve the same relation:

```
SELECT cute AS cuteness,  
       year AS birth  
FROM Dragon;
```

(more readable)



```
SELECT cute AS  
      cuteness, year AS  
      birth FROM Dragon;
```



	cuteness	birth
	10	2010
	-100	2011
	0	2019
	100	2010
	None	2011

Use newlines and whitespace wisely in your SQL queries.
It will simplify your debugging process!

Uniqueness with DISTINCT

To return only unique values, combine SELECT with the DISTINCT keyword

```
SELECT DISTINCT year  
FROM Dragon;
```

Notice that 2010 and 2011 only appear once each in the output.

name	year	cute	year
hiccup	2010	10	2010
drogon	2011	-100	
dragon 2	2019	0	2011
puff	2010	100	2019
smaug	2011	None	

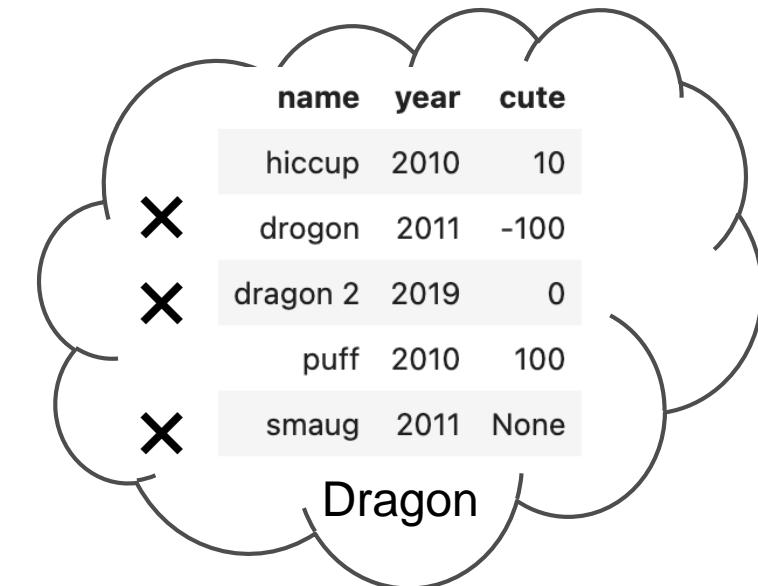
WHERE: Select a rows based on conditions

- To select only some rows of a table, we can use the WHERE keyword.

```
SELECT name, year  
FROM Dragon  
WHERE cute > 0;
```

condition

name	year
hiccup	2010
puff	2010



WHERE: Select a rows based on conditions

- Comparators OR, AND, and NOT let us form more complex conditions.

```
SELECT name, year  
FROM Dragon  
WHERE cute > 0 OR year > 2013;
```

condition

	name	cute	year
1	hiccup	10	2010
2	dragon 2	0	2019
3	puff	100	2010

- Check if values are contained IN a specified list

```
SELECT name, year  
FROM Dragon  
WHERE name IN ('puff', 'hiccup');
```

	name	year
1	hiccup	2010
2	puff	2010

WHERE with NULL Values

NULL (the SQL equivalent of NaN) is stored in a special format – we can't use the “standard” operators =, >, and <.

Instead, check if something IS or IS NOT NULL

```
SELECT name, year  
FROM Dragon  
WHERE year IS NOT NULL;
```

name	cute
hiccup	10
drogon	-100
dragon 2	0
puff	100

Always work with NULLs using the IS operator.
NULL cannot work with standard comparisons:
in fact, NULL = NULL actually returns False!

ORDER BY: Sort rows

- Specify which column(s) we should order the data by

```
SELECT *  
FROM Dragon  
ORDER BY cute DESC;
```

 column



(by default, SQL orders by ascending order: **ASC**)

name	year	cute
puff	2010	100
hiccup	2010	10
dragon 2	2019	0
drogon	2011	-100
smaug	2011	None

ORDER BY: Sort rows

Specify which column(s) we should order the data by

```
SELECT *  
FROM Dragon  
ORDER BY year, cute DESC;
```

Can also order by multiple columns (for tiebreaks)

Sorts year in ascending order and cute in descending order. If you want year to be ordered in descending order as well, you need to specify year DESC, cute DESC;

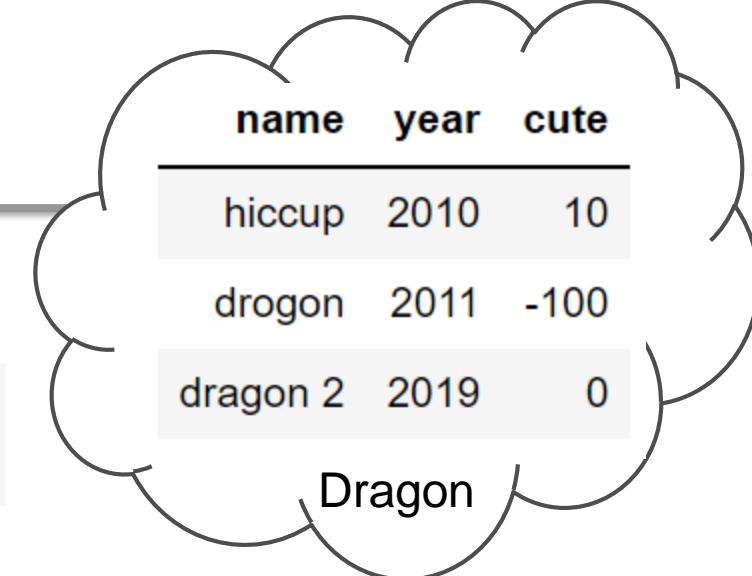
name	year	cute
puff	2010	100
hiccup	2010	10
drogon	2011	-100
smaug	2011	None
dragon 2	2019	0

OFFSET and LIMIT?

1. `SELECT *
FROM Dragon
LIMIT 2;`

A. **name year cute**

hiccup	2010	10
drogon	2011	-100



2. `SELECT *
FROM Dragon
LIMIT 2
OFFSET 1;`

B. **name year cute**

drogon	2011	-100
dragon 2	2019	0

OFFSET and LIMIT

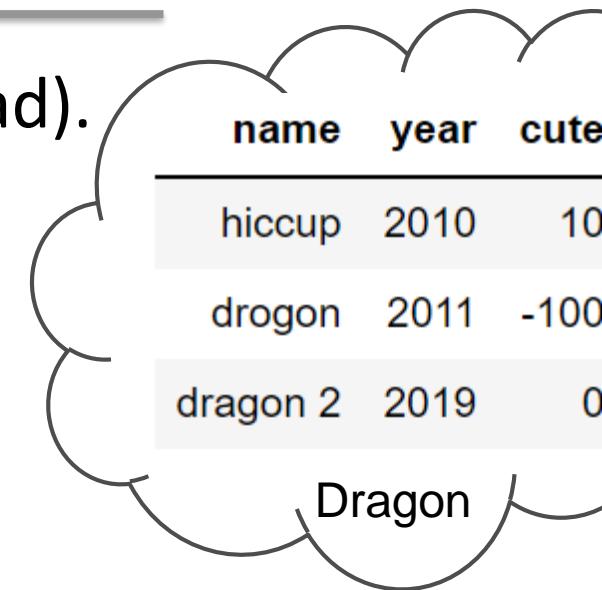
- The LIMIT keyword lets you retrieve N rows (like pandas head).

```
SELECT *  
FROM Dragon  
LIMIT 2;
```

name	year	cute
------	------	------

hiccup	2010	10
--------	------	----

drogon	2011	-100
--------	------	------



- The OFFSET keyword tells SQL to skip the first N rows of the output, then apply LIMIT.

```
SELECT *  
FROM Dragon  
LIMIT 2  
OFFSET 1;
```

name	year	cute
------	------	------

drogon	2011	-100
--------	------	------

dragon 2	2019	0
----------	------	---

⚠️ Unless you use ORDER BY, there is **no guaranteed order** of rows in the relation!

Grouping

- Why Databases?
- Intro to SQL
- Tables and Schema
- Basic Queries
- **Grouping**

The Dish Table

```
SELECT *  
FROM Dish;
```

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

The Dish Table

```
SELECT *  
FROM Dish;
```

Notice the repeated dish types. What if we wanted to investigate trends across each group?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Declarative Programming

Order of operations: SELECT → FROM → WHERE → GROUP BY

```
SELECT type, SUM(cost)  
FROM Dish  
GROUP BY type;
```

Correct! ✓

```
GROUP BY type  
SELECT type, SUM(cost)  
FROM Dish;
```

Incorrect ✗

Always follow the SQL order of operations. Let SQL take care of the rest.

GROUP BY

GROUP BY is similar to pandas groupby().

```
SELECT type  
FROM Dragon  
GROUP BY type;
```

type	name	type	cost
appetizer	ravioli	entree	10
dessert	ramen	entree	13
entree	taco	entree	7
	edamame	appetizer	4
	fries	appetizer	4
	potsticker	appetizer	4
	ice cream	dessert	5

The diagram illustrates the grouping process. It shows a vertical list of categories ('appetizer', 'dessert', 'entree') on the left, with arrows pointing to specific groups of rows in the table on the right. A green arrow points to the first three rows (ravioli, ramen, taco), which are all labeled 'entree'. A purple arrow points to the last two rows (edamame, fries), which are both labeled 'appetizer'. A grey arrow points to the middle row (taco), which is labeled 'entree'.

Aggregating Across Groups

Like pandas, SQL has **aggregate functions**: MAX, SUM, AVG, FIRST, etc.

For more aggregations, see: https://www.sqlite.org/lang_aggfunc.html

```
SELECT type, SUM(cost)
FROM Dish
GROUP BY type;
```

type	SUM(cost)
appetizer	12
dessert	5
entree	30

Wait, something's weird...

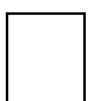
Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```

What do you think will happen?

name	type	cost
ravioli	entree	10
ramen	entree	13
taco	entree	7
edamame	appetizer	4
fries	appetizer	4
potsticker	appetizer	4
ice cream	dessert	5

Dish

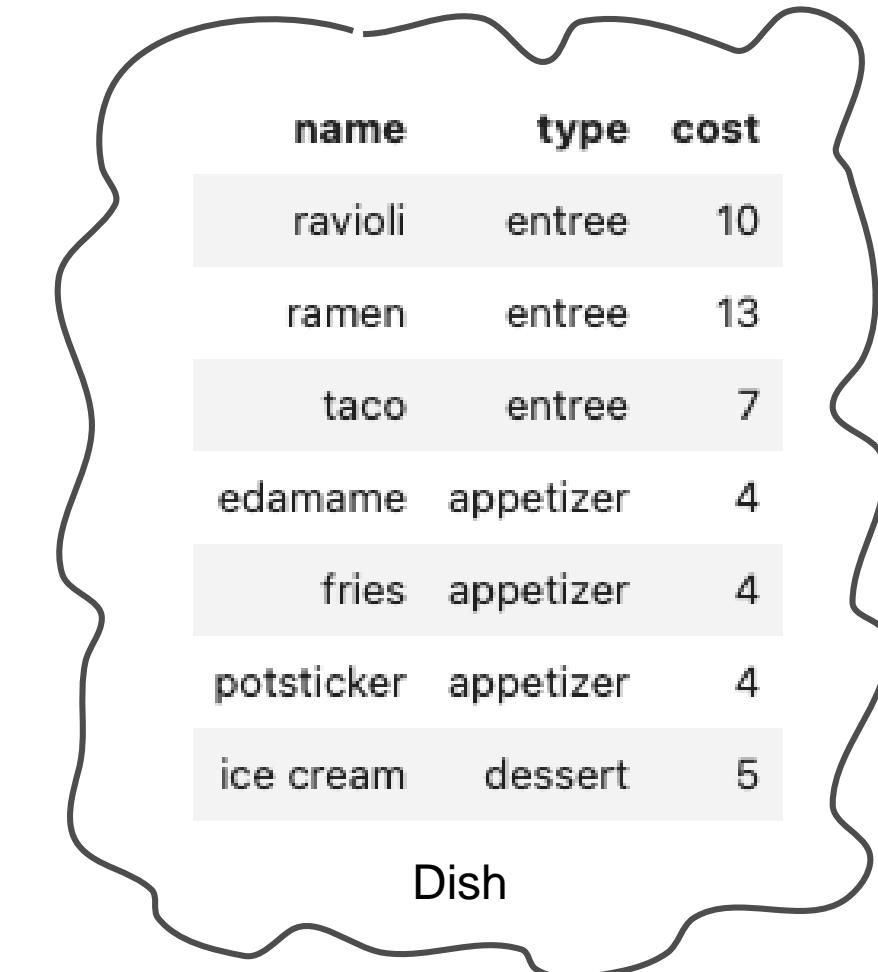


Using Multiple Aggregation Functions

```
SELECT type,  
       SUM(cost),  
       MIN(cost),  
       MAX(name)  
FROM Dish  
GROUP BY type;
```



type	SUM(cost)	MIN(cost)	MAX(name)
appetizer	12	4	potsticker
dessert	5	5	ice cream
entree	30	7	taco



This was much more difficult in pandas!

The COUNT Aggregation

COUNT is used to count the number of rows belonging to a group.

```
SELECT year, COUNT(cute)  
FROM Dragon  
GROUP BY year;
```

Similar to pandas `groupby().count()`

year	COUNT(cute)
2010	2
2011	1
2019	1

COUNT(*) returns the number of rows in each group, including rows with **NULLs**.

```
SELECT year, COUNT(*)  
FROM Dragon  
GROUP BY year;
```

Similar to pandas `groupby().size()`

year	COUNT(*)
2010	2
2011	2
2019	1